

Penalized Generalized Estimating Equations for High-dimensional Longitudinal Data Analysis

Lan Wang

School of Statistics, University of Minnesota, 224 Church Street SE, Minneapolis, MN 55455, U.S.A.

email: wangx346@umn.edu

and

Jianhui Zhou

Department of Statistics, University of Virginia, Charlottesville, VA 22904, USA

email: jz9p@virginia.edu

and

Annie Qu

Department of Statistics, University of Illinois at Urbana-Champaign, Champaign, IL 61820 USA

email: anniequ@illinois.edu

SUMMARY: We consider the penalized generalized estimating equations for analyzing longitudinal data with high-dimensional covariates which often arise in microarray experiments and large-scale health studies. Existing high-dimensional regression procedures often assume independent data and rely on the likelihood function. Construction of a feasible joint likelihood function for high-dimensional longitudinal data is challenging, particularly for correlated discrete outcome data. The penalized generalized estimating equations procedure only requires to specify the first two marginal moments and a working correlation structure. We establish the asymptotic theory in a high-dimensional framework where the number of covariates p_n increases as the number of clusters n increases, and p_n can reach the same order as n . One important feature of the new procedure is that the consistency of model selection holds even if the working correlation structure is misspecified. We evaluate the performance of the proposed method using Monte Carlo simulations and demonstrate its application using a yeast cell-cycle gene expression data set.

000 0000

KEY WORDS: correlated data; diverging number of parameters; GEE; high-dimensional covariates; longitudinal data; marginal regression; variable selection.

1. Introduction

High-dimensional longitudinal data, which consist of repeated measurements on a large number of covariates, have become increasingly common. In many large-scale long-term health studies, such as the well known Framingham Heart Study, many covariates including age, smoking status, cholesterol level, blood pressure were recorded on the participants over the years to describe their health and lifestyles. High-dimensional data arise even more frequently in gene expression experiments. For example, in Section 6 we analyze a yeast cell-cycle gene expression data set where the gene expression measurements were recorded at different time points during the cell cycle. The data set contains 297 cell-cycle regulated genes and the covariates are the binding probabilities for 96 transcription factors. In other examples, even though the number of variables is not large, when we include various interaction effects the total number of covariates in the statistical model can be considerably large. Despite the large number of covariates, it often occurs that only a subset of them are relevant for modelling the response variable. Inclusion of redundant variables may hinder accuracy and efficiency for both estimation and inference. Thus, it is important to develop new statistical methodology and theory of variable selection and estimation for high-dimensional longitudinal data.

The literature on variable selection for longitudinal data is rather limited due to the challenges imposed by incorporating the intra-cluster correlation. Pan (2001) developed a quasi-likelihood information criterion (QIC) which is analogous to AIC; Cantoni et al. (2005) generalized Mallows's C_p criterion; and Wang and Qu (2009) proposed a BIC criterion based on the quadratic inference function. These are best subset type model selection procedures which become computationally intensive when p is moderately large. Fan and Li (2004) and Wang, Li and Huang (2008) studied regularized semiparametric and nonparametric marginal models for continuous responses; Ni, Zhang and Zhang (2009) recently investigated variable

selection for mixed-effects model with continuous responses; Dziak (2006) and Dziak, Li and Qu (2009) discussed the applications of SCAD-penalized quadratic inference function; Xu et al. (2010) considered a GEE based shrinkage estimator with an artificial objective function; Xue, Qu and Zhou (2010) proposed model selection of marginal generalized additive model analysis for correlated data. However, the aforementioned work all assume that the dimension of predictors is fixed. Some of these work only apply to continuous outcome data. For correlated discrete outcome data, the joint likelihood function does not have a close form if the correlation information is taken into account. When the dimension of parameters diverges, numerical approximation to the joint likelihood function tends to be computationally infeasible as it often involves high-dimensional integration.

The generalized estimating equations (GEE) approach is widely applied to longitudinal data analysis (Liang and Zeger, 1986). An important advantage of the GEE approach is that it yields a consistent estimator even if the working correlation structure is misspecified. In this paper, we study the penalized GEE, which solves penalized generalized estimating equations with a non-convex penalty function. Similarly to GEE, the penalized GEE procedure only requires to specify the first two marginal moments and a working correlation matrix. It avoids to specify the full joint likelihood for high-dimensional correlated data, this is particularly appealing for modeling correlated discrete responses.

Johnson, Lin and Zeng (2008) recently derived the asymptotic theory for the penalized estimating equations for independent data. The proposed work contains several important new aspects: we consider multivariate correlated responses and allow the number of covariates to diverge. To establish relevant theory with the diverging number of parameters, we employ rather different techniques than those in Johnson, Lin and Zeng. The proposed work also differs from the earlier penalized GEE approach with the bridge-type of penalty (Fu, 2003), which considers the fixed “ p ” case and mainly aims to solve the collinearity problem.

Under the assumption of a sparse marginal model, we establish the asymptotic theory for the penalized GEE in a high-dimensional framework where the number of candidate covariates p_n increases with the number of clusters n . The new framework allows p_n to be of size comparable to n in the sense that $p_n = O(n)$. The penalized GEE can be effectively solved by an iterative algorithm. Furthermore, we propose a sandwich formula to estimate the asymptotic covariance matrix.

The paper is organized as follows. In Section 2, we introduce GEE and discuss the high-dimensional theory. In Section 3, we propose the penalized GEE and an iterative algorithm. The theory for penalized high-dimensional GEE is presented in Section 4. In Section 5, we report numerical results from Monte-Carlo simulations. Section 6 applies the new method to a yeast cell-cycle gene expression data set. Section 7 provides some discussions.

2. Generalized estimating equations

2.1 Notation and preliminaries

For the j th observation of the i th subject, we observe a response variable Y_{ij} and a p_n -dimensional vector of covariates \mathbf{X}_{ij} , $i = 1, \dots, n$ and $j = 1, \dots, m_i$. The dimension of the covariates p_n is allowed to depend on the number of clusters n . Let $\mathbf{Y}_i = (Y_{i1}, \dots, Y_{im_i})^T$ denote the vector of responses for the i th subject, and let $\mathbf{X}_i = (\mathbf{X}_{i1}, \dots, \mathbf{X}_{im_i})^T$ be the associated $m_i \times p_n$ matrix of covariates. The observations are independent if they are from different subjects, but are correlated if they are from the same subject. In what follows, we assume $m_i = m < \infty$ for simplicity.

We denote the first two marginal moments of Y_{ij} by $\mu_{ij}(\boldsymbol{\beta}_n) := E_{\boldsymbol{\beta}_n}(Y_{ij})$ and $\sigma_{ij}^2(\boldsymbol{\beta}_n) := \text{Var}_{\boldsymbol{\beta}_n}(Y_{ij})$. Let $\boldsymbol{\mu}_i(\boldsymbol{\beta}_n) = (\mu_{i1}(\boldsymbol{\beta}_n), \dots, \mu_{im}(\boldsymbol{\beta}_n))^T$ and $\mathbf{A}_i(\boldsymbol{\beta}_n) = \text{diag}(\sigma_{i1}^2(\boldsymbol{\beta}_n), \dots, \sigma_{im}^2(\boldsymbol{\beta}_n))$. The optimal estimating equation for $\boldsymbol{\beta}_n$ is given by

$$n^{-1} \sum_{i=1}^n \frac{\partial \boldsymbol{\mu}_i(\boldsymbol{\beta}_n)}{\partial \boldsymbol{\beta}_n^T} \mathbf{V}_i^{-1} (\mathbf{Y}_i - \boldsymbol{\mu}_i) = 0, \quad (1)$$

where \mathbf{V}_i is the covariance matrix of \mathbf{Y}_i . In real applications the true intra-cluster covariance structure is often unknown. The GEE procedure adopts a working covariance matrix which is specified through a working correlation matrix $\mathbf{R}(\boldsymbol{\tau})$: $\mathbf{V}_i = \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_n)\mathbf{R}(\boldsymbol{\tau})\mathbf{A}_i^{1/2}(\boldsymbol{\beta}_n)$ where $\boldsymbol{\tau}$ is a finite dimensional parameter. Some commonly used working correlation structures include independence, AR-1, equally correlated (also called compound symmetry) or unstructured correlation, among others. For a given working correlation structure, $\boldsymbol{\tau}$ can be estimated using residual-based moment method.

As in Liang and Zeger (1986), we assume that the marginal density of Y_{ij} comes from a canonical exponential family. Then we can write $\mu_{ij}(\boldsymbol{\beta}_n) = a(\theta_{ij})$ and $\sigma_{ij}^2(\boldsymbol{\beta}_n) = \phi \dot{a}(\theta_{ij})$, where $\theta_{ij} = \mathbf{X}_{ij}^T \boldsymbol{\beta}_n$, for a differentiable function $a(\cdot)$ and a scaling constant ϕ . With the estimated working correlation matrix $\hat{\mathbf{R}}$, the estimating equations in (1) reduce to

$$n^{-1} \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_n) \hat{\mathbf{R}}^{-1} \mathbf{A}_i^{-1/2}(\boldsymbol{\beta}_n) (\mathbf{Y}_i - \boldsymbol{\mu}_i(\boldsymbol{\beta}_n)) = 0. \quad (2)$$

We formally define the GEE estimator as the solution $\hat{\boldsymbol{\beta}}_n$ of the above estimating equations. For eas of exposition, we assume $\phi = 1$ in the rest of the paper. Xie and Yang (2003) and Balan and Schiopu-Kratina (2005) established rigorous large sample theory for the GEE when the number of clusters n is large but the dimension of covariates p is fixed.

2.2 GEE with diverging number of covariates

In a “diverging p ” asymptotic framework but without the sparsity assumption, Wang (2011) studied the consistency, asymptotic normality and validity of the sandwich variance formula of GEE. This framework allows one to adopt a growingly more complex statistical model in order to reduce the modeling bias.

For a vector or a matrix \mathbf{B} , denote $\|\mathbf{B}\| = [\text{Tr}(\mathbf{B}\mathbf{B}^T)]^{1/2}$ as its Frobenius norm, which reduces to the Euclidean norm in the vector case. It can be shown that a preliminary GEE estimator $\tilde{\boldsymbol{\beta}}_n$ obtained by solving the independence working estimating equation $\sum_{i=1}^n \mathbf{X}_i^T (\mathbf{Y}_i -$

$\mu_i(\beta_n) = 0$ satisfies $\|\tilde{\beta}_n - \beta_{n0}\| = O_p(\sqrt{p_n/n})$ when $p_n \rightarrow \infty$ and $p_n^2 n^{-1} = o(1)$, where β_{n0} is the true value of β_n . Under rather general conditions, using the initial consistent estimator $\tilde{\beta}_n$, the residual-based estimated working correlation matrix $\hat{\mathbf{R}}$ satisfies $\|\hat{\mathbf{R}}^{-1} - \bar{\mathbf{R}}^{-1}\| = O_p(\sqrt{p_n/n})$, where $\bar{\mathbf{R}}$ is a constant positive-definite matrix but not necessarily the true correlation matrix \mathbf{R}_0 .

Let

$$\begin{aligned}\bar{\mathbf{M}}_n(\beta_n) &= \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\beta_n) \bar{\mathbf{R}}^{-1} \mathbf{R}_0 \bar{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\beta_n) \mathbf{X}_i, \\ \bar{\mathbf{H}}_n(\beta_n) &= \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\beta_n) \bar{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\beta_n) \mathbf{X}_i.\end{aligned}$$

Under general regularity conditions, Wang (2011) showed that if $p_n^2 n^{-1} = o(1)$, then the generalized estimating equation (2) has a root $\hat{\beta}_n$ such that $\|\hat{\beta}_n - \beta_{n0}\| = O_p(\sqrt{p_n/n})$. Furthermore, if $p_n^3 n^{-1} = o(1)$, then $\forall \alpha_n \in R^{p_n}$ such that $\|\alpha_n\| = 1$,

$$\alpha_n^T \bar{\mathbf{M}}_n^{-1/2}(\beta_{n0}) \bar{\mathbf{H}}_n(\beta_{n0}) (\hat{\beta}_n - \beta_{n0}) \rightarrow N(0, 1) \quad \text{in distribution.}$$

For the purpose of variable selection for high-dimensional longitudinal data, we assume from now on that the marginal mean regression model is sparse in the sense that most of the covariates have zero coefficients. The above results provide some insight on how sparse the true underlying model can be such that we are still confident about the statistical inference after dimension reduction by variable selection. Here we allow the true underlying model size (i.e., number of nonzero coefficients) to increase with the number of clusters but require its dimension to be of order $o(n^{1/3})$.

3. Penalized generalized estimating equations

3.1 Methodology

We consider the penalized generalized estimating equations for simultaneous estimation and variable selection. Specifically, the penalized generalized estimating functions are defined as

$$\mathbf{U}_n(\boldsymbol{\beta}_n) = \mathbf{S}_n(\boldsymbol{\beta}_n) - \mathbf{q}_{\lambda_n}(|\boldsymbol{\beta}_n|)\text{sign}(\boldsymbol{\beta}_n), \quad (3)$$

where

$$\mathbf{S}_n(\boldsymbol{\beta}_n) = n^{-1} \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_n) \hat{\mathbf{R}}^{-1} \mathbf{A}_i^{-1/2}(\boldsymbol{\beta}_n) (\mathbf{Y}_i - \boldsymbol{\mu}_i(\boldsymbol{\beta}_n)) \quad (4)$$

are the estimating functions defining the GEE, $\mathbf{q}_{\lambda_n}(|\boldsymbol{\beta}_n|) = (q_{\lambda_n}(|\boldsymbol{\beta}_{n1}|), \dots, q_{\lambda_n}(|\boldsymbol{\beta}_{np_n}|))^T$ is a p_n -dimensional vector of penalty functions, and $\text{sign}(\boldsymbol{\beta}_n) = (\text{sign}(\boldsymbol{\beta}_{n1}), \dots, \text{sign}(\boldsymbol{\beta}_{np_n}))^T$ with $\text{sign}(t) = I(t > 0) - I(t < 0)$. The notation $\mathbf{q}_{\lambda_n}(|\boldsymbol{\beta}_n|)\text{sign}(\boldsymbol{\beta}_n)$ denotes the component-wise product. The tuning parameter λ_n determines the amount of shrinkage.

Since $U_n(\boldsymbol{\beta})$ has discontinuous points, an exact solution to $U_n(\boldsymbol{\beta}) = 0$ may not exist. We formally define the penalized GEE estimator $\hat{\boldsymbol{\beta}}_n$ to be an approximate solution: $\mathbf{U}_n(\hat{\boldsymbol{\beta}}_n) = o(a_n)$ for a sequence $a_n \rightarrow 0$. The rate of a_n will be made clearer in Theorem 1 below. Alternatively, we may define the penalized GEE estimator as an asymptotic zero crossing, see REMARK 2 below.

To gain insights into the penalized GEE, we first note that the penalty function $q_{\lambda_n}(|\beta_j|)$ is zero for a large value of $|\beta_j|$ and is relatively large for a small value of $|\beta_j|$. As a result, the generalized estimating function $\mathbf{S}_{nj}(\boldsymbol{\beta}_n)$, the j th component of $\mathbf{S}_n(\boldsymbol{\beta}_n)$, is not penalized if β_{nj} is large in magnitude; while the penalty level is high if β_{nj} is close (but not equal) to zero and therefore forces its estimator to be shrunk to zero. Once an estimated coefficient is shrunk to zero, it is excluded from the final selected model.

Although different penalty functions can be potentially adopted, in this paper we consider

the non-convex SCAD penalty which is given by

$$q_{\lambda_n}(\theta) = \lambda_n \left\{ I(\theta \leq \lambda_n) + \frac{(a\lambda_n - \theta)_+}{(a-1)\lambda_n} I(\theta > \lambda_n) \right\}$$

for $\theta \geq 0$ and some $a > 2$. In the context of penalized likelihood, Fan and Li (2001) demonstrated that the SCAD penalty simultaneously achieves three desirable properties of variable selection: unbiasedness, sparsity and continuity. In contrast, the Lasso penalty (L_1 penalty) does not satisfy the unbiasedness condition, the L_q penalty with $q > 1$ does not satisfy the sparsity condition, and the L_q penalty with $0 \leq q < 1$ does not satisfy the continuity condition. Following the suggestion of Fan and Li, we take $a = 3.7$.

3.2 Algorithm

To solve the penalized GEE, similarly as in Johnson, Lin and Zeng (2008), we use an iterative algorithm which combines the minorization-maximization (MM) algorithm for non-convex penalty of Hunter and Li (2005) with the Newton-Raphson algorithm for the GEE.

For a small $\epsilon > 0$, the MM algorithm suggests that the penalized GEE estimator $\hat{\boldsymbol{\beta}}_n = (\hat{\beta}_{n1}, \dots, \hat{\beta}_{np_n})^T$ approximately satisfies

$$S_{nj}(\hat{\boldsymbol{\beta}}_n) - nq_{\lambda_n}(|\hat{\beta}_{nj}|)\text{sign}(\hat{\beta}_{nj})\frac{|\hat{\beta}_{nj}|}{\epsilon + |\hat{\beta}_{nj}|} = 0, \quad j = 1, \dots, p_n, \quad (5)$$

where $S_{nj}(\hat{\boldsymbol{\beta}}_n)$ denotes the j th element of $\mathbf{S}_n(\hat{\boldsymbol{\beta}}_n)$. In the numerical analysis, we take ϵ to be a fixed small number 10^{-6} .

Applying the Newton-Raphson algorithm of GEE to $S_{nj}(\hat{\boldsymbol{\beta}}_n)$ and combining with (5), we obtain the following iterative algorithm for solving the penalized GEE:

$$\hat{\boldsymbol{\beta}}_n^k = \hat{\boldsymbol{\beta}}_n^{k-1} + [\mathbf{H}_n(\hat{\boldsymbol{\beta}}_n^{k-1}) + n\mathbf{E}_n(\hat{\boldsymbol{\beta}}_n^{k-1})]^{-1}[\mathbf{S}_n(\hat{\boldsymbol{\beta}}_n^{k-1}) - n\mathbf{E}_n(\hat{\boldsymbol{\beta}}_n^{k-1})\hat{\boldsymbol{\beta}}_n^{k-1}],$$

where

$$\begin{aligned}\mathbf{H}_n(\hat{\boldsymbol{\beta}}_n^{k-1}) &= \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\hat{\boldsymbol{\beta}}_n^{k-1}) \hat{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\hat{\boldsymbol{\beta}}_n^{k-1}) \mathbf{X}_i, \\ \mathbf{E}_n(\hat{\boldsymbol{\beta}}_n^{k-1}) &= \text{diag}\left\{ \frac{q_{\lambda_n}(|\hat{\beta}_{n1}|+)}{\epsilon + |\hat{\beta}_{n1}|}, \dots, \frac{q_{\lambda_n}(|\hat{\beta}_{np_n}|+)}{\epsilon + |\hat{\beta}_{np_n}|} \right\}.\end{aligned}$$

Given a selected tuning parameter, we repeat the above algorithm to update $\hat{\boldsymbol{\beta}}_n^k$ for 30 iterations with the initial value $\hat{\boldsymbol{\beta}}_n^0 = (0, 0, \dots, 0)^T$. Our numerical experience in Section 5 indicates that the criterion $\sum_{j=1}^{p_n} |\hat{\boldsymbol{\beta}}_j^{k+1} - \hat{\boldsymbol{\beta}}_j^k| < 0.0001$ is usually met within 30 iterations. In practice, we select the tuning parameter λ_n using cross-validation. More specifically, we randomly split the data into several non-overlapping subsets of approximately equal size. We remove one subset and fit the model to the remaining data, and estimate the prediction error (the loss function is the negative loglikelihood under working independence assumption) from the removed observations. This is repeated for each subset, the estimated prediction errors are aggregated, and the best tuning parameter is selected by minimizing the aggregated estimated prediction error over a fine grid.

Extension of the proposed penalized GEE estimator to the cases of unequal m_i is straightforward. We vary the dimension of A_i and replace \hat{R} by \hat{R}_i , which is the $m_i \times m_i$ matrix using the specified working correlation structure and the corresponding initial parameter $\boldsymbol{\tau}$ estimator.

4. Asymptotic theory for high-dimensional penalized GEE

In this section, we present the theory of the penalized GEE in a “large n , diverging p ” framework, which allows p_n to be of similar order as n (see REMARK 3 below). We denote the true value of $\boldsymbol{\beta}_{n0}$ by $\boldsymbol{\beta}_{n0} = (\boldsymbol{\beta}_{n10}^T, \boldsymbol{\beta}_{n20}^T)^T$. The covariates matrix is partitioned into $\mathbf{X}_i = (\mathbf{X}_{i1}, \mathbf{X}_{i2})$ accordingly. Without loss of generality, it is assumed that $\boldsymbol{\beta}_{n20} = \mathbf{0}$ and that the elements of $\boldsymbol{\beta}_{n10}$ are all nonzero. We also denote the dimension of $\boldsymbol{\beta}_{n10}$ by s_n , where

s_n may be fixed or grow with n . Following the discussions in Section 2.2, we assume that $s_n = o(n^{1/3})$.

For the asymptotic theory, we require the following regularity conditions, which are assumed for technical convenience and may be further relaxed.

(A1) X_{ij} , $1 \leq i \leq n$, $1 \leq j \leq m$, are uniformly bounded.

(A2) The unknown parameter β_n belongs to a compact subset $\mathcal{B} \subseteq R^{p_n}$ and the true parameter value β_{n0} lies in the interior of \mathcal{B} .

(A3) There exist finite positive constants b_1 and b_2 such that

$$b_1 \leq \lambda_{\min} \left(n^{-1} \sum_{i=1}^n \mathbf{X}_i^T \mathbf{X}_i \right) \leq \lambda_{\max} \left(n^{-1} \sum_{i=1}^n \mathbf{X}_i^T \mathbf{X}_i \right) \leq b_2,$$

where λ_{\min} (or λ_{\max}) denotes the minimum (or maximum) eigenvalue of a matrix.

(A4) The common true correlation matrix \mathbf{R}_0 has eigenvalues bounded away from zero and $+\infty$. The estimated working correlation matrix $\hat{\mathbf{R}}$ satisfies $\|\hat{\mathbf{R}}^{-1} - \bar{\mathbf{R}}^{-1}\| = O_p(\sqrt{s_n/n})$, where $\bar{\mathbf{R}}$ is a constant positive-definite matrix with eigenvalues bounded away from zero and $+\infty$.

(A5) Let $\epsilon_i(\beta_n) = (\epsilon_{i1}(\beta_n), \dots, \epsilon_{im}(\beta_n))^T = \mathbf{A}_i^{-1/2}(\beta_n)(\mathbf{Y}_i - \boldsymbol{\mu}_i(\beta_n))$. There exists a finite constant $M_1 > 0$ such that $E(\|\epsilon_i(\beta_{n0})\|^{2+\delta}) \leq M_1$, for all i and some $\delta > 0$; and there exist positive constants M_2 and M_3 such that $E[\exp(M_2|\epsilon_{ij}(\beta_{n0})|)|\mathbf{X}_i] \leq M_3$, uniformly in $i = 1, \dots, n$, $j = 1, \dots, m$.

(A6) Let $B_n = \{\beta_n : \|\beta_n - \beta_{n0}\| \leq \Delta\sqrt{p_n/n}\}$, then $\dot{\mu}(\mathbf{X}_{ij}^T \beta_n)$, $1 \leq i \leq n$, $1 \leq j \leq m$, are uniformly bounded away from 0 and ∞ on B_n ; $\ddot{\mu}(\mathbf{X}_{ij}^T \beta_n)$ and $\mu^{(3)}(\mathbf{X}_{ij}^T \beta_n)$, $1 \leq i \leq n$, $1 \leq j \leq m$, are uniformly bounded by a finite positive constant M_2 on B_n .

(A7) Assuming $\min_{1 \leq j \leq s_n} |\beta_{n0j}|/\lambda_n \rightarrow \infty$ as $n \rightarrow \infty$ and $s_n^3 n^{-1} = o(1)$, $\lambda_n \rightarrow 0$, $s_n^2 (\log n)^4 = o(n\lambda_n^2)$, $\log(p_n) = o(n\lambda_n^2/(\log n)^2)$, $p_n s_n^4 (\log n)^6 = o(n^2 \lambda_n^2)$ and $p_n s_n^3 (\log n)^8 = o(n^2 \lambda_n^4)$.

REMARK 1. The first part of condition (A5) is similar to the condition in Lemma 2 of Xie and Yang (2003) and condition (\tilde{N}_δ) in Balan and Schiopu-Kratina (2005); the second part is satisfied for Gaussian distribution, sub-Gaussian distribution and Poisson distribution, etc. Condition (A6) requires $\mu_{ij}^{(k)}(\mathbf{X}_{ij}^T \boldsymbol{\beta}_n)$, which denotes the k th derivative of $\mu_{ij}(t)$ evaluated at $\mathbf{X}_{ij}^T \boldsymbol{\beta}_n$, to be uniformly bounded when $\boldsymbol{\beta}_n$ is in a local neighborhood around $\boldsymbol{\beta}_{n0}$, $k = 1, 2, 3$. This condition is generally satisfied for the GEE. For example, when the marginal model follows a Poisson distribution, $\mu(t) = \exp(t)$, thus $\mu_{ij}^{(k)}(\mathbf{X}_{ij}^T \boldsymbol{\beta}_n) = \exp(\mathbf{X}_{ij}^T \boldsymbol{\beta}_n)$, $k = 1, 2, 3$, are uniformly bounded around $\boldsymbol{\beta}_{n0}$ on B_n .

The regularized procedures based on minimizing the sum of an objective function and a penalty term have been well studied for analyzing high-dimensional data, for example LASSO (Tibshirani 1996), SCAD (Fan and Li, 2001; Fan and Peng, 2004) and MCP (Zhang, 2010), among others. However, for the estimating equation approach, alternative techniques are necessary to establish the asymptotic theory. Theorem 1 below characterizes the asymptotic properties of the penalized GEE procedure.

THEOREM 1: *Assuming conditions (A1)-(A7), there exists an approximate penalized GEE solution $\hat{\boldsymbol{\beta}}_n = (\hat{\boldsymbol{\beta}}_{n1}^T, \hat{\boldsymbol{\beta}}_{n2}^T)^T$ which satisfies the following properties*

(1)

$$P(|U_{nj}(\hat{\boldsymbol{\beta}}_n)| = 0, \quad j = 1, \dots, s_n) \rightarrow 1 \quad (6)$$

$$P\left(|U_{nj}(\hat{\boldsymbol{\beta}}_n)| \leq \frac{\lambda_n}{\log n}, \quad j = s_n + 1, \dots, p_n\right) \rightarrow 1 \quad (7)$$

(2) $P(\hat{\boldsymbol{\beta}}_{n2} = \mathbf{0}) \rightarrow 1$;

(3) $\forall \boldsymbol{\alpha}_n \in R^{s_n}$ such that $\|\boldsymbol{\alpha}_n\| = 1$,

$$\boldsymbol{\alpha}_n^T \overline{\mathbf{M}}_{n1}^{-1/2}(\boldsymbol{\beta}_{n0}) \overline{\mathbf{H}}_{n1}(\boldsymbol{\beta}_{n0}) (\hat{\boldsymbol{\beta}}_{n1} - \boldsymbol{\beta}_{n10}) \rightarrow N(0, 1) \quad \text{in distribution,}$$

where

$$\begin{aligned}\overline{\mathbf{M}}_{n1}(\boldsymbol{\beta}_{n0}) &= \sum_{i=1}^n \mathbf{X}_{i1}^T \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_{n0}) \overline{\mathbf{R}}^{-1} \mathbf{R}_0 \overline{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_{n0}) \mathbf{X}_{i1}, \\ \overline{\mathbf{H}}_{n1}(\boldsymbol{\beta}_{n0}) &= \sum_{i=1}^n \mathbf{X}_{i1}^T \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_{n0}) \overline{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\boldsymbol{\beta}_{n0}) \mathbf{X}_{i1}.\end{aligned}$$

Properties (2) and (3) in Theorem 1 are often referred to as the oracle property of variable selection, that is, the procedure estimates the true zero-coefficient as zero with probability approaching one and estimates the nonzero coefficients as efficiently as if the true model is known in advance. Property (1) provides a more precise characterization for the approximate solution of the penalized GEE.

REMARK 2. Alternatively, we may define the penalized GEE estimator as an approximate zero crossing as in Johnson, Lin and Zeng (2008). One can show that an estimator $\widehat{\boldsymbol{\beta}}_n$ in Theorem 1 is also an approximate zero crossing in the sense that a small perturbation of $\widehat{\boldsymbol{\beta}}_n$ changes the sign of the penalized estimating equations. To understand this more clearly, we consider a perturbation of the j th component of $\widehat{\boldsymbol{\beta}}_n$, for some $j = s_1 + 1, \dots, p_n$. Such a small perturbation leads to a penalty term of order λ_n which dominates the GEE part (of order $\lambda_n(\log n)^{-1}$). Therefore, the direction of the perturbation determines the sign of the penalized GEE function.

REMARK 3. The number of candidate covariates p_n allowed in the asymptotic theory depends on the dimension of the true model. When the true underlying model has a fixed dimension, which is often assumed in the literature, the above theory allows $p_n = o(n^\alpha)$ with $0 < \alpha < \frac{4}{3}$. Note that for such p_n condition (A7) is satisfied for some $\lambda_n = O(n^{-\gamma})$, where $0 < \gamma < (2 - \alpha)/4$, as long as $\min_{1 \leq j \leq s_n} |\beta_{n0j}|/\lambda_n \rightarrow \infty$. Thus the theory allows p_n to be of similar size as n in the sense that $p_n = O(n)$.

From the algorithm in Section 3.2, we obtain the following sandwich formula to estimate

the asymptotic covariance matrix of $\hat{\beta}_n$:

$$\text{Cov}(\hat{\beta}_n) \approx [\mathbf{H}_n(\hat{\beta}_n) + n\mathbf{E}_n(\hat{\beta}_n)]^{-1} \mathbf{M}_n(\hat{\beta}_n) [\mathbf{H}_n(\hat{\beta}_n) + n\mathbf{E}_n(\hat{\beta}_n)]^{-1},$$

where \mathbf{H}_n and \mathbf{E}_n are defined in Section 3.2, and

$$\mathbf{M}_n(\hat{\beta}_n) = \sum_{i=1}^n \mathbf{X}_i^T \mathbf{A}_i^{1/2}(\hat{\beta}_n) \hat{\mathbf{R}}^{-1} [\boldsymbol{\epsilon}_i(\hat{\beta}_n) \boldsymbol{\epsilon}_i^T(\hat{\beta}_n)] \hat{\mathbf{R}}^{-1} \mathbf{A}_i^{1/2}(\hat{\beta}_n) \mathbf{X}_i,$$

with $\boldsymbol{\epsilon}_i(\hat{\beta}_n) = \mathbf{A}_i^{-1/2}(\hat{\beta}_n)(\mathbf{Y}_i - \boldsymbol{\mu}_i(\hat{\beta}_n))$.

5. Monte Carlo simulations

We next conduct simulations to evaluate the performance of the proposed penalized GEE procedure for both normal and binary responses. We compare it with the unpenalized GEE and the oracle GEE (i.e., the GEE with the true marginal regression model known in advance). For each procedure, we consider three different working correlation structures: independence, exchangeable and AR(1). For each setup in the simulations, we generate 100 data sets and apply the iterative algorithm in Section 3.2 to estimate β_{n0} , where we run 30 iterations for each simulated data set. We select the tuning parameter λ_n in the SCAD penalty function using a 4-fold cross validation. At the end of the iteration, if an estimated coefficient has magnitude below the cut-off value 10^{-3} , it is considered as zero.

For evaluating estimation accuracy, we report the estimated *mean squared error* (MSE), defined as $\sum_{j=1}^{100} \|\hat{\beta}_n^j - \beta_{n0}\|^2 / 100$, where $\hat{\beta}_n^j$ is the estimator of β_{n0} obtained using the j th generated data set. To evaluate model selection performance, we report the proportion of times the methods under-selecting (U), over-selecting (O) and exactly selecting (EXACT) the covariates with nonzero coefficients. We also report the average true positives (denoted by TP, the average number of selected covariates that correspond to the nonzero coefficients in the underlying model) and the average false positives (denoted by FP, the average number of selected covariates that correspond to the zero coefficients in the underlying model).

Example 1. The correlated normal responses are generated from the model

$$Y_{ij} = \mathbf{X}_{ij}^T \boldsymbol{\beta} + \epsilon_{ij},$$

where $i = 1, \dots, 200$, $j = 1, \dots, 4$, $\mathbf{X}_{ij} = (x_{ij,1}, \dots, x_{ij,200})^T$ is a vector of 200 covariates and $\boldsymbol{\beta} = (2.0, 3.0, 1.5, 2.0, 0, \dots, 0)^T$. For the covariates, we generate $x_{ij,1}$ from the Bernoulli(0.5) distribution, and $x_{ij,2}$ to $x_{ij,200}$ from the multivariate normal distribution with mean 0 and an AR(1) covariance matrix with marginal variance 1 and auto-correlation coefficient 0.5. The random errors $(\epsilon_{i1}, \dots, \epsilon_{i4})^T$ are generated from the multivariate normal distribution with marginal mean 0, marginal variance 1 and an exchangeable correlation matrix with parameter ρ . We consider $\rho = 0.5$ and 0.8 to represent different strength of within cluster correlation.

[Table 1 about here.]

The results of Table 1 summarize the estimation accuracy and model selection properties of the penalized GEE, the unpenalized GEE and the oracle GEE for three different working correlation matrices and two different values of ρ . We observe that in terms of estimation accuracy the penalized GEE procedure performs closely to the oracle GEE, and significantly reduces the MSE of the unpenalized GEE estimator. Using the true correlation structure (exchangeable) in penalized GEE gives the smallest MSE, with greater gain when the within cluster association is stronger. Furthermore, we observe that the unpenalized GEE generally does not lead to a sparse model. The penalized GEE successfully selects all covariates with nonzero coefficients and has a fairly small number of false positives. When the true working correlation structure is used, the probability of identifying the exact underlying model (i.e., both false positive and false negative are zero) is about 70%.

To further investigate the performance of the penalized GEE estimator, Table 2 reports its bias, the estimated standard deviation (calculated from the sandwich variance formula) and

the empirical standard deviation for estimating β_i , $i = 1, \dots, 4$. We also use the sandwich variance formula to construct approximate 95% confidence intervals for the four nonzero coefficients based on the asymptotic normality theory and report the empirical coverage probabilities. The estimated standard deviation is close to the empirical standard deviation; and the empirical coverage probability is close to 95%. This indicated good performance of the sandwich variance formula.

[Table 2 about here.]

Example 2. The correlated binary responses have marginal mean π_{ij} satisfying

$$\log \frac{\pi_{ij}}{1 - \pi_{ij}} = \mathbf{X}_{ij}^T \boldsymbol{\beta},$$

where $i = 1, \dots, 400$, $j = 1, \dots, 10$, $\mathbf{X}_{ij} = (x_{ij,1}, \dots, x_{ij,50})^T$ is a vector of 50 covariates, and $\boldsymbol{\beta} = (0.7, -0.7, -0.4, 0, 0, \dots, 0)^T$. For the covariates, we generate $x_{ij,k}$ independently such that $x_{ij,k} \sim \text{Uniform}(0, 1)$. We use the R package “mvtBinaryEP” to generate the correlated binary responses with an exchangeable correlation structure ($\rho = 0.4$) within each cluster.

We note that Example 2 is more challenging than Example 1 in two aspects: first, the response is binary thus contains much less information than a continuous response; secondly, we include a small signal ($\beta_3 = -0.4$) in the nonzero coefficients. Table 3 summarizes the results for estimation accuracy and model selection properties of the penalized GEE, the unpenalized GEE and the oracle GEE for three different working correlation matrices. We observe that the penalized GEE significantly improves the estimation accuracy of the unpenalized GEE. The penalized GEE has high true positive rate and low false positive rate for variable selection. We also observe that when the true correlation structure is used, the PGEE has high probability (97%) to catch this small signal ($\beta_3 = -0.4$); but when the

independence working correlation structure is used, there is a significantly higher chance to miss this small signal.

[Table 3 about here.]

Similarly as Table 2, Table 4 reports the bias, the estimated standard deviation (calculated from the sandwich variance formula), the empirical standard deviation, and the empirical coverage probability of 95% confidence interval for estimating β_i , $i = 1, 2, 3$, when the penalized GEE procedure is applied. When the true correlation structure is used, the empirical coverage probabilities for β_1 and β_2 are close to 95% (the estimated standard deviation is also close to the empirical standard deviation); the empirical coverage probability for the small coefficient β_3 is close to 90% (the estimated standard deviation is smaller comparing to the empirical standard deviation). However, when the independence working correlation structure is used, the confidence intervals undercover the true value. These observations suggest that modeling covariance structure is important when some of the signals are relatively weak.

[Table 4 about here.]

6. Yeast cell-cycle gene expression data analysis

We apply the penalized GEE to the yeast cell-cycle gene expression data collected in the CDC15 experiment of Spellman et al. (1998). The experiment recorded genome-wide mRNA levels for 6178 yeast ORFs (abbreviation for open reading frames, which are DNA sequences that can determine which amino acids will be encoded by a gene) at 7-minute intervals for 119 minutes, which covers two cell-cycle periods for a total of 18 time points.

The cell cycle is a tightly regulated life process where cells grow, replicate their DNA, segregate their chromosomes and divide into as many daughter cells as the environment allows. The cell-cycle process is commonly divided into M/G1-G1-S-G2-M stages, where the M stage stands for mitosis during which nuclear (chromosome separation) and cytoplasmic

(cytokinesis) and division occur, the G1 stage stands for GAP 1; the S stage stands for synthesis during which DNA replication occurs; the G2 stage stands for GAP 2. Spellman et al.'s experiment identified approximately 800 genes which vary in a periodic fashion during the yeast cell cycle, however little was known about the regulation of most of these genes. Transcription factors (TFs) have been observed to play critical roles in gene expression regulation. A transcription factor (sometimes called a sequence-specific DNA-binding factor) is a protein that binds to specific DNA sequences, thereby controlling the flow (or transcription) of genetic information from DNA to mRNA. We next apply the penalized GEE to identify the TFs that influence the gene expression level at each stage of the cell process. This is important to understanding how the cell cycle is regulated and how cell cycles regulate other biological processes.

Similarly as in Luan and Li (2003) and Wang, Li and Huang (2008), we analyze a subset of 297 cell-cycle-regularized genes. The response variable Y_{ij} is the log-transformed gene expression level of gene i measured at time point j ; the covariates x_{ik} , $k = 1, \dots, 96$, is the matching score of the binding probability of the k th transcription factor on the promoter region of the i th gene. The binding probability is computed using a mixture modeling approach based on data from a ChIP binding experiment, see Wang, Chen and Li (2007) for details. We apply the penalized GEE to each of the five stages (each stage contains a few time points from the cycle) of the cell-cycle process using the following model

$$Y_{ij} = \alpha_0 + \alpha_1 t_{ij} + \sum_{k=1}^{96} \beta_k x_{ik} + \epsilon_{ij},$$

where x_{ik} , $k = 1, \dots, 96$, is standardized to have mean zero and variance 1; t_{ij} denotes time. We impose penalties on the β_k 's. Table 5 summarizes the number of TFs identified when three different working correlation structures for ϵ_{ij} are adopted: independence, AR-1 and exchangeable.

[Table 5 about here.]

Our analysis reveals that at each of the five stages the selected TFs, in terms of both numbers and the specific TFs, are not sensitive to the choice of the working correlation structure. Due to the space limitation, we only report the number of selected TFs in Table 5. Some of these selected TFs have already been confirmed by biological experiments using genome-wide binding method. For examples, Fkh1, Fkh2 and Mcm1 are three TFs that have been proved important for stage G2 in the aforementioned biological experiments and they have been selected by the penalized GEE for stage G2. Furthermore, the sets of TFs selected at different stages have only small overlaps. This suggests that different TFs play important roles at different stages of the cell-cycle process, which has also been observed by the biologists.

7. Discussions

In general, when the number of covariates is large, identifying the exact underlying model is a challenging task, in particular when some of the nonzero signals are relatively weak. From a practical point of view, it is often satisfactory to identify a model that includes all important variables along with a small number of false positives. In general, under fitting model is more serious than over fitting in model selection. Although oracle property is suggested by Theorem 1, the practical performance of the penalized GEE may be influenced by two factors: (1) the multiple roots problem of the estimating equations approach and (2) the tuning parameter selection. For the later problem, it has been observed that cross-validation tends to select an over-fitted model. A high-dimensional BIC approach (Chen and Chen, 2008) might be more effective for selecting λ_n , but it is challenging to establish the relevant theory in the GEE setting. This will be a topic of future research.

Acknowledgements

The authors would like to thank the two referees, the AE and the Co-Editor for their careful reading and constructive comments which substantially improved an earlier draft. Wang's research is supported by National Science Foundation grant DMS-1007603; Zhou's research is supported by National Science Foundation grant DMS-0906665 and Qu's research is supported by National Science Foundation grant DMS-0906660.

Supplementary Materials

The proof of Theorem 1 is available under the Paper Information link at the the Biometrics website <http://www.tibs.org/biometrics>.

References

- Balan, R. M. and Schiopu-Kratina, I. (2005). Asymptotic results with generalized estimating equations for longitudinal data. *Annals of Statistic*, **32**, 522-541.
- Cantoni, E., Flemming, J. M. and Ronchetti, E. (2005). Variable selection for marginal longitudinal generalized linear models. *Biometrics*, **61**, 507-514.
- Chen, J. and Chen, Z. (2008). Extended Bayesian information criterion for model selection with large model space. *Biometrika*, **95**, 759-771.
- Dziak, J. J. (2006). Penalized quadratic inference function for variable selection in longitudinal research. Ph.D. Thesis, the Pennsylvania State University.
- Dziak, J. J., Li, R and Qu, A. (2009) An overview on quadratic inference function approaches for longitudinal data. In *Frontiers of Statistics, Vol 1: New developments in biostatistics and bioinformatics*, edited by Fan, J., Liu, J. S. and Lin, X. World scientific publishing, Chapter 3, 49-72.
- Fan, J. and Li, R. (2001). Variable selection via nonconcave penalized likelihood and its oracle properties. *Journal of the American Statistical Association*, **96**, 1348-1360.

- Fan, J. and Li, R. (2004). New estimation and model selection procedures for semiparametric modeling in longitudinal data analysis. *Journal of the American Statistical Association*, **99**, 710-723.
- Fan, J. and Peng, H. (2004). Nonconcave penalized likelihood with a diverging number of parameters. *Annals of Statistics*, **32**, 928-961.
- Fu, W. J. (2003). Penalized estimating equations. *Biometrics*, **35**, 109-148.
- Hunter, D. R. and Li, R. (2005). Variable selection using MM algorithms. *Annals of Statistics*, **33**, 1617-1642.
- Johnson, B., Lin, D. Y. and Zeng, D. (2008). Penalized estimating functions and variable selection in semiparametric regression models. *Journal of the American Statistical Association*, **103**, 672-680.
- Liang, K. Y. and Zeger, S. L. (1986). Longitudinal data analysis using generalized linear models. *Biometrika* **73**, 12-22.
- Luan, Y., and Li, H. (2003). Clustering of time-course gene expression data using a mixed-effects model with Bsplines. *Bioinformatics*, **19**, 474-482.
- Pan, W. (2001). Akaike's information criterion in generalized estimating equations. *Biometrics*, **57**, 120-125.
- Spellman, P. T., Sherlock, G., Zhang, M. Q., Iyer, V. R., Anders, K., Eisen, M. B., Brown, P. O., Botstein, D., and Futcher, B. (1998). Comprehensive identification of cell cycle-regulated genes of the yeast *saccharomyces cerevisiae* by microarray hybridization. *Molecular Biology of Cell*, **9**, 3273-3297.
- Tibshirani, R. J. (1996). Regression shrinkage and selection via the Lasso. *Journal of the Royal Statistical Society, Series B*, **58**, 267-288.
- Wang, L. (2011). GEE analysis of clustered binary data with diverging number of covariates. *Annals of Statistics*, **39**, 389-417.

- Wang, L., Chen, G. and Li, H. (2007). Group SCAD regression analysis for microarray time course gene expression data. *Bioinformatics*, **23**, 1486-1494.
- Wang L, Li, H. and Huang. J (2008). Variable selection in nonparametric varying- coefficient models for analysis of repeated measurements. *Journal of the American Statistical Association*, **103**, 1556-1569.
- Wang, L. and Qu, A. (2009). Consistent model selection and data-driven smooth tests for longitudinal data in the estimating equations approach. *Journal of the Royal Statistical Society, Series B*, **71**, 177-190.
- Xiao, N., Zhang, D., and Zhang H. H. (2009). Variable selection for semiparametric mixed models in longitudinal studies. To appear in *Biometrics*.
- Xie, M. and Yang, Y. (2003). Asymptotics for generalized estimating equations with large cluster sizes. *The Annals of Statistics*, **31**, 310-347.
- Xu, P., Wu, P., Wang, Y. and Zhu, L.X. (2010). A GEE based shrinkage estimation for the generalized linear model in longitudinal data analysis. Technical report, Department of Mathematics, Hong Kong Baptist University, Hong Kong.
- Xue, L, Qu, A. and Zhou, J. (2010). Efficient estimation and model selection for the marginal generalized additive model for correlated data. To appear in *Journal of the American Statistical Association*.
- Zhang, C. H. (2010). Nearly unbiased variable selection under minimax concave penalty. *Annals of Statistics*, **38**, 894-942.

Table 1: Correlated continuous responses ($n = 200$, $p_n = 200$): comparison of GEE, oracle GEE and penalized GEE with three different working correlation matrices (independence, exchangeable and ar1).

	MSE	U	O	EXACT	TP	FP
$\rho = 0.5$						
GEE.indep	0.568	0.00	1.00	0.00	4.00	193.02
GEE.exch	0.381	0.00	1.00	0.00	4.00	192.45
GEE.ar1	0.458	0.00	1.00	0.00	4.00	192.66
Oracle.indep	0.009	-	-	-	-	-
Oracle.exch	0.006	-	-	-	-	-
Oracle.ar1	0.007	-	-	-	-	-
PGEE.indep	0.009	0.00	0.85	0.15	4.00	2.02
PGEE.exch	0.008	0.00	0.33	0.67	4.00	3.30
PGEE.ar1	0.008	0.00	0.38	0.62	4.00	3.00
$\rho = 0.8$						
GEE.indep	0.568	0.00	1.00	0.00	4.00	193.01
GEE.exch	0.165	0.00	1.00	0.00	4.00	190.44
GEE.ar1	0.211	0.00	1.00	0.00	4.00	191.53
Oracle.indep	0.010	-	-	-	-	-
Oracle.exch	0.003	-	-	-	-	-
Oracle.ar1	0.003	-	-	-	-	-
PGEE.indep	0.011	0.00	0.83	0.17	4.00	2.15
PGEE.exch	0.004	0.00	0.33	0.67	4.00	4.23
PGEE.ar1	0.005	0.00	0.35	0.65	4.00	4.02

Table 2: Correlated continuous responses ($n = 200$, $p_n = 200$): performance of the penalized GEE for three working correlation matrices (independence, exchangeable and ar1). Bias denotes the absolute value of the empirical bias; SD1 denotes the estimated standard deviation using the sandwich variance estimator; SD2 denotes the sample standard deviation; CP denotes the empirical coverage probability of the 95% confidence interval.

Correlation		β_1	β_2	β_3	β_4
$\rho = 0.5$					
indep	Bias	0.002	0.000	0.002	0.003
indep	SD1	0.067	0.041	0.046	0.041
indep	SD2	0.065	0.038	0.043	0.041
indep	CP	96.00	96.00	97.00	97.00
exch	Bias	0.000	0.002	0.001	0.000
exch	SD1	0.051	0.032	0.036	0.032
exch	SD2	0.053	0.030	0.036	0.031
exch	CP	95.00	96.00	93.00	97.00
ar1	Bias	0.001	0.001	0.001	0.000
ar1	SD1	0.054	0.034	0.038	0.034
ar1	SD2	0.056	0.032	0.038	0.034
ar1	CP	95.00	97.00	95.00	95.00
$\rho = 0.8$					
indep	Bias	0.002	0.001	0.002	0.005
indep	SD1	0.074	0.041	0.046	0.040
indep	SD2	0.072	0.038	0.043	0.042
indep	CP	97.00	96.00	97.00	97.00
exch	Bias	0.000	0.001	0.001	0.000
exch	SD1	0.037	0.021	0.024	0.021
exch	SD2	0.039	0.020	0.024	0.020
exch	CP	95.00	94.00	94.00	97.00
ar1	Bias	0.002	0.000	0.000	0.000
ar1	SD1	0.042	0.023	0.026	0.023
ar1	SD2	0.043	0.022	0.026	0.023
ar1	CP	96.00	96.00	95.00	96.00

Table 3: Correlated binary responses ($n = 400$, $p_n = 50$): comparison of GEE, oracle GEE and penalized GEE with three different working correlation matrices (independence, exchangeable and ar1).

	MSE	U	O	EXACT	TP	FP
GEE.indep	0.635	0.00	1.00	0.00	3.00	46.51
GEE.exch	0.421	0.00	1.00	0.00	3.00	46.60
GEE.ar1	0.576	0.00	1.00	0.00	3.00	46.59
Oracle.indep	0.027	-	-	-	-	-
Oracle.exch	0.018	-	-	-	-	-
Oracle.ar1	0.025	-	-	-	-	-
PGEE.indep	0.111	0.28	0.32	0.40	2.72	0.93
PGEE.exch	0.049	0.03	0.37	0.60	2.97	1.05
PGEE.ar1	0.081	0.09	0.37	0.54	2.91	1.33

Table 4: Correlated binary responses ($n = 400$, $p_n = 50$): performance of the penalized GEE for three working correlation matrices (independence, exchangeable and ar1). Bias denotes the absolute value of the empirical bias; SD1 denotes the estimated standard deviation using the sandwich variance estimator; SD2 denotes the sample standard deviation; CP denotes the empirical coverage probability of the 95% confidence interval.

Correlation		β_1	β_2	β_3
indep	Bias	0.038	0.048	0.107
indep	SD1	0.099	0.101	0.066
indep	SD2	0.119	0.125	0.221
indep	CP	85.00	87.00	64.00
exch	Bias	0.001	0.005	0.027
exch	SD1	0.081	0.083	0.076
exch	SD2	0.078	0.082	0.128
exch	CP	94.00	95.00	88.00
ar1	Bias	0.010	0.017	0.050
ar1	SD1	0.089	0.090	0.077
ar1	SD2	0.087	0.095	0.171
ar1	CP	96.00	98.00	81.00

Table 5: Number of TFs selected for each stage in the yeast cell-cycle process with the penalized GEE procedure

Correlation	M/G1	G1	S	G2	M
indep	20	19	19	10	22
exch	20	18	18	10	18
ar1	23	17	18	10	19